

# Thực hành trên lop-answer

December 20, 2021

## 1 Thực hành Transformers

Trong bài này, ta sẽ thực hành cài đặt Transformer

### 1.0.1 1. Cài đặt và import thư viện

```
[ ]: !which python3
```

```
[ ]: !pip3 install spacy dill  
!pip3 install torchtext  
!pip3 install pandas
```

```
[ ]: !python3 -m spacy download en && python3 -m spacy download fr
```

```
[ ]: import torch.nn as nn  
import torch  
import torchtext  
import copy  
import math  
import torch.nn.functional as F
```

```
[ ]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

### 1.0.2 2. Cài đặt từng module của Transformer

```
[ ]: class Embedder(nn.Module):  
    def __init__(self, vocab_size, dim):  
        super().__init__()  
        self.embed = nn.Embedding(vocab_size, dim)  
  
    def forward(self, x):  
        return self.embed(x)
```

Position Embedding Class:

```
[ ]: # Positional encoding  
class PositionalEncoder(nn.Module):  
    def __init__(self, dim, max_seq_len=300):
```

```

    super().__init__()
    self.dim = dim

    # create a constant 'pe' matrix with values dependant on
    # pos and i
    pe = torch.zeros(max_seq_len, dim)
    for pos in range(max_seq_len):
        for i in range(0, dim, 2):
            pe[pos, i] = math.sin(pos / (10000 ** ((2*i)/dim)))
            pe[pos, i+1] = math.cos(pos / (10000 ** ((2* (i+1))/dim)))

    pe = pe.unsqueeze(0)
    self.register_buffer('pe', pe)

    def forward(self, x):
        # make embeddings relatively larger
        x = x * math.sqrt(self.dim)
        # add constant to embedding
        seq_len = x.size(1)
        x = x + Variable(self.pe[:, :seq_len], requires_grad=False).to(device)
        return x

```

**Multi Head Attention:** We first start with implementing attention function

Attention of  $q$

```

[ ]: def attention(q, k, v, d_k, mask=None, dropout=None):
    scores = torch.matmul(q, k.transpose(-2, -1)) / math.sqrt(d_k)
    if mask is not None:
        mask = mask.unsqueeze(1)
        scores = scores.masked_fill(mask == 0, -1e9)

    scores = F.softmax(scores, dim=-1)

    if dropout is not None:
        scores = dropout(scores)

    output = torch.matmul(scores, v)
    return output

```

```

[ ]: # Multi-headed attention
class MultiHeadAttention(nn.Module):
    def __init__(self, heads, dim, dropout=0.1):
        super().__init__()
        self.dim = dim
        self.dim_head = dim//heads
        self.h = heads
        self.q_linear = nn.Linear(dim, dim)

```

```

self.k_linear = nn.Linear(dim, dim)
self.v_linear = nn.Linear(dim, dim)
self.dropout = nn.Dropout(dropout)
self.out = nn.Linear(dim, dim)

def forward(self, q, k, v, mask=None):
    bs = q.size(0)
    # perform linear operation and split into h heads
    k = self.k_linear(k).view(bs, -1, self.h, self.dim_head)
    q = self.q_linear(q).view(bs, -1, self.h, self.dim_head)
    v = self.v_linear(v).view(bs, -1, self.h, self.dim_head)
    # transpose to get dimensions bs * h * sl * dim
    k = k.transpose(1, 2)
    q = q.transpose(1, 2)
    v = v.transpose(1, 2)
    # calculate attention using the function we will define next
    scores = attention(q, k, v, self.dim, mask, self.dropout)
    # concatenate heads and put through final linear layer
    concat = scores.transpose(1,2).contiguous().view(bs, -1, self.dim)
    output = self.out(concat)
    return output

```

```

[ ]: class FeedForward(nn.Module):
    def __init__(self, d_model, d_ff=2048, dropout = 0.1):
        super().__init__()
        # We set d_ff as a default to 2048
        self.linear_1 = nn.Linear(d_model, d_ff)
        self.dropout = nn.Dropout(dropout)
        self.linear_2 = nn.Linear(d_ff, d_model)
    def forward(self, x):
        x = self.dropout(F.relu(self.linear_1(x)))
        x = self.linear_2(x)
        return x

```

```

[ ]: class Norm(nn.Module):
    def __init__(self, d_model, eps = 1e-6):
        super().__init__()

        self.size = d_model
        # create two learnable parameters to calibrate normalisation
        self.alpha = nn.Parameter(torch.ones(self.size))
        self.bias = nn.Parameter(torch.zeros(self.size))
        self.eps = eps
    def forward(self, x):
        norm = self.alpha * (x - x.mean(dim=-1, keepdim=True)) \
            / (x.std(dim=-1, keepdim=True) + self.eps) + self.bias
        return norm

```

```
[ ]: # build an encoder layer with one multi-head attention layer and one  
# feed-forward layer
```

```
class EncoderLayer(nn.Module):  
    def __init__(self, d_model, heads, dropout = 0.1):  
        super().__init__()  
        self.norm_1 = Norm(d_model)  
        self.norm_2 = Norm(d_model)  
        self.attn = MultiHeadAttention(heads, d_model)  
        self.ff = FeedForward(d_model)  
        self.dropout_1 = nn.Dropout(dropout)  
        self.dropout_2 = nn.Dropout(dropout)  
  
    def forward(self, x, mask):  
        x2 = self.norm_1(x)  
        x = x + self.dropout_1(self.attn(x2,x2,x2,mask))  
        x2 = self.norm_2(x)  
        x = x + self.dropout_2(self.ff(x2))  
        return x
```

```
[ ]: # build a decoder layer with two multi-head attention layers and  
# one feed-forward layer
```

```
class DecoderLayer(nn.Module):  
    def __init__(self, d_model, heads, dropout=0.1):  
        super().__init__()  
        self.norm_1 = Norm(d_model)  
        self.norm_2 = Norm(d_model)  
        self.norm_3 = Norm(d_model)  
  
        self.dropout_1 = nn.Dropout(dropout)  
        self.dropout_2 = nn.Dropout(dropout)  
        self.dropout_3 = nn.Dropout(dropout)  
  
        self.attn_1 = MultiHeadAttention(heads, d_model)  
        self.attn_2 = MultiHeadAttention(heads, d_model)  
        self.ff = FeedForward(d_model).cuda()  
  
    def forward(self, x, e_outputs, src_mask, trg_mask):  
        x2 = self.norm_1(x)  
        x = x + self.dropout_1(self.attn_1(x2, x2, x2, trg_mask))  
        x2 = self.norm_2(x)  
        x = x + self.dropout_2(self.attn_2(x2, e_outputs, e_outputs,  
src_mask))  
        x2 = self.norm_3(x)  
        x = x + self.dropout_3(self.ff(x2))  
        return x
```

```
def get_clones(module, N):
    return nn.ModuleList([copy.deepcopy(module) for i in range(N)])
```

```
[ ]: class Encoder(nn.Module):
    def __init__(self, vocab_size, d_model, N, heads):
        super().__init__()
        self.N = N
        self.embed = Embedder(vocab_size, d_model)
        self.pe = PositionalEncoder(d_model)
        self.layers = get_clones(EncoderLayer(d_model, heads), N)
        self.norm = Norm(d_model)

    def forward(self, src, mask):
        x = self.embed(src)
        x = self.pe(x)
        for i in range(self.N):
            x = self.layers[i](x, mask)
        return self.norm(x)

class Decoder(nn.Module):
    def __init__(self, vocab_size, d_model, N, heads):
        super().__init__()
        self.N = N
        self.embed = Embedder(vocab_size, d_model)
        self.pe = PositionalEncoder(d_model)
        self.layers = get_clones(DecoderLayer(d_model, heads), N)
        self.norm = Norm(d_model)
    def forward(self, trg, e_outputs, src_mask, trg_mask):
        x = self.embed(trg)
        x = self.pe(x)
        for i in range(self.N):
            x = self.layers[i](x, e_outputs, src_mask, trg_mask)
        return self.norm(x)
```

```
[ ]: class Transformer(nn.Module):
    def __init__(self, src_vocab, trg_vocab, d_model, N, heads):
        super().__init__()
        self.encoder = Encoder(src_vocab, d_model, N, heads)
        self.decoder = Decoder(trg_vocab, d_model, N, heads)
        self.out = nn.Linear(d_model, trg_vocab)
    def forward(self, src, trg, src_mask, trg_mask):
        e_outputs = self.encoder(src, src_mask)
        d_output = self.decoder(trg, e_outputs, src_mask, trg_mask)
        output = self.out(d_output)
        return output # we don't perform softmax on the output as this will be
        ↪ handled
# automatically by our loss function
```

### 1.0.3 3. Chuẩn bị và tiền xử lý dữ liệu

```
[ ]: import spacy
import re

# Tokenize

class tokenize(object):

    def __init__(self, lang):
        self.nlp = spacy.load(lang)

    def tokenizer(self, sentence):
        sentence = re.sub(
            r"[\*\"'\"'\n\\...\\+\\-\\/\\=\\(\\)‘•: \\[\\]\\\\|'\\!;]", " ", str(sentence))
        sentence = re.sub(r"[ ]+", " ", sentence)
        sentence = re.sub(r"!+", "!", sentence)
        sentence = re.sub(r"\\,", ",", sentence)
        sentence = re.sub(r"\\?", "?", sentence)
        sentence = sentence.lower()
        return [tok.text for tok in self.nlp.tokenizer(sentence) if tok.text != '↵']
```

```
[ ]: # Creating batch
from torchtext.legacy import data
import numpy as np
from torch.autograd import Variable

def nopeak_mask(size, opt):
    np_mask = np.triu(np.ones((1, size, size)),
        k=1).astype('uint8')
    np_mask = Variable(torch.from_numpy(np_mask) == 0)
    np_mask = np_mask.to(device)
    return np_mask

def create_masks(src, trg, opt):

    src_mask = (src != opt.src_pad).unsqueeze(-2)

    if trg is not None:
        trg.to(device)
        trg_mask = (trg != opt.trg_pad).unsqueeze(-2).to(device)
        size = trg.size(1) # get seq_len for matrix
        np_mask = nopeak_mask(size, opt)
        trg_mask = trg_mask & np_mask
```

```

else:
    trg_mask = None
    return src_mask, trg_mask

# patch on Torchtext's batching process that makes it more efficient
# from http://nlp.seas.harvard.edu/2018/04/03/attention.
# →html#position-wise-feed-forward-networks

class MyIterator(data.Iterator):
    def create_batches(self):
        if self.train:
            def pool(d, random_shuffler):
                for p in data.batch(d, self.batch_size * 100):
                    p_batch = data.batch(
                        sorted(p, key=self.sort_key),
                        self.batch_size, self.batch_size_fn)
                    for b in random_shuffler(list(p_batch)):
                        yield b
            self.batches = pool(self.data(), self.random_shuffler)

        else:
            self.batches = []
            for b in data.batch(self.data(), self.batch_size,
                               self.batch_size_fn):
                self.batches.append(sorted(b, key=self.sort_key))

global max_src_in_batch, max_tgt_in_batch

def batch_size_fn(new, count, sofar):
    "Keep augmenting batch and calculate total number of tokens + padding."
    global max_src_in_batch, max_tgt_in_batch
    if count == 1:
        max_src_in_batch = 0
        max_tgt_in_batch = 0
    max_src_in_batch = max(max_src_in_batch, len(new.src))
    max_tgt_in_batch = max(max_tgt_in_batch, len(new.trg) + 2)
    src_elements = count * max_src_in_batch
    tgt_elements = count * max_tgt_in_batch
    return max(src_elements, tgt_elements)

```

```

[ ]: import pandas as pd
import torchtext
from torchtext.legacy import data
import os
import dill as pickle

def read_data(opt):

```

```

if opt.src_data is not None:
    try:
        opt.src_data = open(opt.src_data).read().strip().split('\n')
    except:
        print("error: " + opt.src_data + " file not found")
        quit()

if opt.trg_data is not None:
    try:
        opt.trg_data = open(opt.trg_data).read().strip().split('\n')
    except:
        print("error: " + opt.trg_data + " file not found")
        quit()

def create_fields(opt):
    spacy_langs = ['en', 'fr', 'de', 'es', 'pt', 'it', 'nl']
    src_lang = opt.src_lang[0:2]
    trg_lang = opt.trg_lang[0:2]
    if src_lang not in spacy_langs:
        print('invalid src language: ' + opt.src_lang + 'supported languages :_
↪' + spacy_langs)
    if trg_lang not in spacy_langs:
        print('invalid trg language: ' + opt.trg_lang + 'supported languages :_
↪' + spacy_langs)

    print("loading spacy tokenizers...")

    t_src = tokenize(opt.src_lang)
    t_trg = tokenize(opt.trg_lang)
    TRG = data.Field(lower=True, tokenize=t_trg.tokenizer, init_token='<sos>',_
↪eos_token='<eos>')
    SRC = data.Field(lower=True, tokenize=t_src.tokenizer)

    return(SRC, TRG)

def create_dataset(opt, SRC, TRG):

    print("creating dataset and iterator... ")

    raw_data = {'src' : [line for line in opt.src_data], 'trg': [line for line_
↪in opt.trg_data]}
    df = pd.DataFrame(raw_data, columns=["src", "trg"])

    mask = (df['src'].str.count(' ') < opt.max_strlen) & (df['trg'].str.count('_
↪') < opt.max_strlen)
    df = df.loc[mask]

```



```

df.to_csv("translate_transformer_temp.csv", index=False)

data_fields = [('src', SRC), ('trg', TRG)]
train = data.TabularDataset('./translate_transformer_temp.csv',
↪format='csv', fields=data_fields)

train_iter = MyIterator(train, batch_size=opt.batchsize, device=device,
                        repeat=False, sort_key=lambda x: (len(x.src), len(x.
↪trg)),
                        batch_size_fn=batch_size_fn, train=True, shuffle=True)

os.remove('translate_transformer_temp.csv')
SRC.build_vocab(train)
TRG.build_vocab(train)
opt.src_pad = SRC.vocab.stoi['<pad>']
opt.trg_pad = TRG.vocab.stoi['<pad>']

opt.train_len = get_len(train_iter)

return train_iter

def get_len(train):

    for i, b in enumerate(train):
        pass

    return i

```

#### 1.0.4 4. Cài đặt giải thuật tối ưu và huấn luyện mô hình

```

[ ]: # Optimizer
class CosineWithRestarts(torch.optim.lr_scheduler._LRScheduler):
    """
    Cosine annealing with restarts.
    Parameters
    -----
    optimizer : torch.optim.Optimizer
    T_max : int
        The maximum number of iterations within the first cycle.
    eta_min : float, optional (default: 0)
        The minimum learning rate.
    last_epoch : int, optional (default: -1)
        The index of the last epoch.
    """

    def __init__(self,
                 optimizer: torch.optim.Optimizer,

```

```

        T_max: int,
        eta_min: float = 0.,
        last_epoch: int = -1,
        factor: float = 1.) -> None:
# pylint: disable=invalid-name
self.T_max = T_max
self.eta_min = eta_min
self.factor = factor
self._last_restart: int = 0
self._cycle_counter: int = 0
self._cycle_factor: float = 1.
self._updated_cycle_len: int = T_max
self._initialized: bool = False
super(CosineWithRestarts, self).__init__(optimizer, last_epoch)

def get_lr(self):
    """Get updated learning rate."""
    # HACK: We need to check if this is the first time get_lr() was called,
    ↪ since
    # we want to start with step = 0, but _LRScheduler calls get_lr with
    # last_epoch + 1 when initialized.
    if not self._initialized:
        self._initialized = True
        return self.base_lrs

    step = self.last_epoch + 1
    self._cycle_counter = step - self._last_restart

    lrs = [
        (
            self.eta_min + ((lr - self.eta_min) / 2) *
            (
                np.cos(
                    np.pi *
                    ((self._cycle_counter) % self._updated_cycle_len) /
                    self._updated_cycle_len
                ) + 1
            )
        ) for lr in self.base_lrs
    ]

    if self._cycle_counter % self._updated_cycle_len == 0:
        # Adjust the cycle length.
        self._cycle_factor *= self.factor
        self._cycle_counter = 0
        self._updated_cycle_len = int(self._cycle_factor * self.T_max)
        self._last_restart = step

```

```
return lrs
```

```
[ ]: !mkdir data
!wget https://raw.githubusercontent.com/SamLynnEvans/Transformer/master/data/
↪english.txt
!mv english.txt data
!wget https://raw.githubusercontent.com/SamLynnEvans/Transformer/master/data/
↪french.txt data/french.txt
!mv french.txt data
```

```
[ ]: def get_model(opt, src_vocab, trg_vocab):

    assert opt.d_model % opt.heads == 0
    assert opt.dropout < 1

    model = Transformer(src_vocab, trg_vocab, opt.d_model, opt.n_layers, opt.
↪heads)

    if opt.load_weights is not None:
        print("loading pretrained weights...")
        model.load_state_dict(torch.load(f'{opt.load_weights}/model_weights'))
    else:
        for p in model.parameters():
            if p.dim() > 1:
                nn.init.xavier_uniform_(p)

    if opt.device == 0:
        model = model.cuda()

    return model
```

```
[ ]: import time

def train_model(model, opt):

    print("training model...")
    model.train()
    start = time.time()
    if opt.checkpoint > 0:
        cptime = time.time()

    for epoch in range(opt.epochs):

        total_loss = 0
        print("    %dm: epoch %d [%s]  %d%%  loss = %s" %\
```

```

        ((time.time() - start)//60, epoch + 1, "".join(' '*20), 0, '...'),
        end='\r')

    if opt.checkpoint > 0:
        torch.save(model.state_dict(), 'weights/model_weights')

    for i, batch in enumerate(opt.train):

        src = batch.src.transpose(0,1).to(device)
        trg = batch.trg.transpose(0,1).to(device)
        trg_input = trg[:, :-1].to(device)
        src_mask, trg_mask = create_masks(src, trg_input, opt)
        preds = model(src, trg_input, src_mask, trg_mask)
        ys = trg[:, 1:].contiguous().view(-1)
        opt.optimizer.zero_grad()
        loss = F.cross_entropy(preds.view(-1, preds.size(-1)), ys,
        ignore_index=opt.trg_pad)
        loss.backward()
        opt.optimizer.step()

        total_loss += loss.item()

    if (i + 1) % opt.printevery == 0:
        p = int(100 * (i + 1) / opt.train_len)
        avg_loss = total_loss/opt.printevery
        print("    %dm: epoch %d [%s%s]  %d%%  loss = %.3f" %\
              ((time.time() - start)//60, epoch + 1, "".join('#'*(p//5)),
        "".join(' '*20), p, avg_loss))
        total_loss = 0

    if opt.checkpoint > 0 and ((time.time()-cptime)//60) // opt.
    checkpoint >= 1:
        torch.save(model.state_dict(), 'weights/model_weights')
        cptime = time.time()

    print("%dm: epoch %d [%s%s]  %d%%  loss = %.3f\nepoch %d complete, loss
    = %.03f" %\
          ((time.time() - start)//60, epoch + 1, "".join('#'*(100//5)), "".join('
    '*20), 100, avg_loss, epoch + 1, avg_loss))

class Opt(object):
    pass

def main():
    opt = Opt()

```

```

opt.src_data = "data/english.txt"
opt.trg_data = "data/french.txt"
opt.src_lang = "en_core_web_sm"
opt.trg_lang = 'fr_core_news_sm'
opt.epochs = 2
opt.d_model=512
opt.n_layers=6
opt.heads=8
opt.dropout=0.1
opt.batchsize=1500
opt.printevery=100
opt.lr=0.0001
opt.max_strlen=80
opt.checkpoint = 0
opt.no_cuda = False
opt.load_weights = None

opt.device = 0
if opt.device == 0:
    assert torch.cuda.is_available()

read_data(opt)
SRC, TRG = create_fields(opt)
opt.train = create_dataset(opt, SRC, TRG)
model = get_model(opt, len(SRC.vocab), len(TRG.vocab)).to(device)

opt.optimizer = torch.optim.Adam(model.parameters(), lr=opt.lr, betas=(0.9,
→0.98), eps=1e-9)

if opt.checkpoint > 0:
    print("model weights will be saved every %d minutes and at end of epoch
→to directory weights/"%(opt.checkpoint))

train_model(model, opt)

# for asking about further training use while true loop, and return
if __name__ == "__main__":
    main()

```