

TrainingDNNs2_pytorch_answer

December 20, 2021

1 Training Deep Neural Networks (Tiếp)

Trong buổi này, chúng ta tiếp tục làm quen với một số kỹ thuật huấn luyện mạng nơ-ron

```
[ ]: !nvidia-smi
# from google.colab import drive
# drive.mount('/content/drive')

import torch
import torch.nn as nn
import torch.optim as optim
import numpy as np
import glob
import cv2
import torch.nn.functional as F
from torch.autograd import Variable

import torchvision
import torchvision.transforms as transforms

from torch.nn import CrossEntropyLoss, Dropout, Softmax, Linear, Conv2d,
↳ LayerNorm
import matplotlib.pyplot as plt
from torchsummary import summary
```

Fri Sep 10 08:45:59 2021

```
+-----+
| NVIDIA-SMI 470.63.01      Driver Version: 460.32.03      CUDA Version: 11.2      |
+-----+-----+-----+
| GPU  Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp   Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                       |                    |     MIG M.     |
+=====+=====+=====+
|   0   Tesla K80           Off   | 00000000:00:04:0 Off  |             0      |
| N/A   69C    P8      33W / 149W |      0MiB / 11441MiB |      0%      Default  |
|                                       |                    |             N/A     |
+-----+-----+-----+
```

Processes:						
GPU	GI	CI	PID	Type	Process name	GPU Memory
	ID	ID				Usage
=====						
No running processes found						

1.0.1 1. Cài đặt BatchNorm

```
[ ]: def compare_bn(bn1, bn2):
    err = False
    if not torch.allclose(bn1.running_mean, bn2.running_mean):
        print('Diff in running_mean: {} vs {}'.format(
            bn1.running_mean, bn2.running_mean))
        err = True

    if not torch.allclose(bn1.running_var, bn2.running_var):
        print('Diff in running_var: {} vs {}'.format(
            bn1.running_var, bn2.running_var))
        err = True

    if bn1.affine and bn2.affine:
        if not torch.allclose(bn1.weight, bn2.weight):
            print('Diff in weight: {} vs {}'.format(
                bn1.weight, bn2.weight))
            err = True

        if not torch.allclose(bn1.bias, bn2.bias):
            print('Diff in bias: {} vs {}'.format(
                bn1.bias, bn2.bias))
            err = True

    if not err:
        print('All parameters are equal!')

class MyBatchNorm2d(nn.BatchNorm2d):
    def __init__(self, num_features, eps=1e-5, momentum=0.1,
                 affine=True, track_running_stats=True):
        super(MyBatchNorm2d, self).__init__(
            num_features, eps, momentum, affine, track_running_stats)

    def forward(self, input):
        self._check_input_dim(input)

        exponential_average_factor = 0.0
```

```

        if self.training and self.track_running_stats:
            if self.num_batches_tracked is not None:
                self.num_batches_tracked += 1
                if self.momentum is None: # use cumulative moving average
                    exponential_average_factor = 1.0 / float(self.
→num_batches_tracked)
                else: # use exponential moving average
                    exponential_average_factor = self.momentum

        # calculate running estimates
        if self.training:
            mean = input.mean([0, 2, 3])
            # use biased var in train
            var = input.var([0, 2, 3], unbiased=False)
            n = input.numel() / input.size(1)
            with torch.no_grad():
                self.running_mean = exponential_average_factor * mean\
                    + (1 - exponential_average_factor) * self.running_mean
                # update running_var with unbiased var
                self.running_var = exponential_average_factor * var * n / (n -
→1)\
                    + (1 - exponential_average_factor) * self.running_var
            else:
                mean = self.running_mean
                var = self.running_var

            input = (input - mean[None, :, None, None]) / (torch.sqrt(var[None, :,
→None, None] + self.eps))
            if self.affine:
                input = input * self.weight[None, :, None, None] + self.bias[None, :
→, None, None]

        return input

# Init BatchNorm layers
my_bn = MyBatchNorm2d(3, affine=True)
bn = nn.BatchNorm2d(3, affine=True)

compare_bn(my_bn, bn) # weight and bias should be different
# Load weight and bias
my_bn.load_state_dict(bn.state_dict())
compare_bn(my_bn, bn)

# Run train
for _ in range(10):

```

```

scale = torch.randint(1, 10, (1,)).float()
bias = torch.randint(-10, 10, (1,)).float()
x = torch.randn(10, 3, 100, 100) * scale + bias
out1 = my_bn(x)
out2 = bn(x)
compare_bn(my_bn, bn)

torch.allclose(out1, out2)
print('Max diff: ', (out1 - out2).abs().max())

# Run eval
my_bn.eval()
bn.eval()
for _ in range(10):
    scale = torch.randint(1, 10, (1,)).float()
    bias = torch.randint(-10, 10, (1,)).float()
    x = torch.randn(10, 3, 100, 100) * scale + bias
    out1 = my_bn(x)
    out2 = bn(x)
    compare_bn(my_bn, bn)

torch.allclose(out1, out2)

```

```

All parameters are equal!
All parameters are equal!
All parameters are equal!
Max diff:  tensor(4.7684e-07, grad_fn=<MaxBackward1>)
All parameters are equal!
Max diff:  tensor(4.7684e-07, grad_fn=<MaxBackward1>)
All parameters are equal!
Max diff:  tensor(4.7684e-07, grad_fn=<MaxBackward1>)
All parameters are equal!
Max diff:  tensor(4.7684e-07, grad_fn=<MaxBackward1>)
All parameters are equal!
Max diff:  tensor(4.7684e-07, grad_fn=<MaxBackward1>)
All parameters are equal!
Max diff:  tensor(4.7684e-07, grad_fn=<MaxBackward1>)
All parameters are equal!
Max diff:  tensor(4.7684e-07, grad_fn=<MaxBackward1>)
All parameters are equal!
Max diff:  tensor(4.7684e-07, grad_fn=<MaxBackward1>)
All parameters are equal!
Max diff:  tensor(4.7684e-07, grad_fn=<MaxBackward1>)
All parameters are equal!
Max diff:  tensor(4.7684e-07, grad_fn=<MaxBackward1>)
All parameters are equal!
All parameters are equal!
All parameters are equal!

```

All parameters are equal!
All parameters are equal!
All parameters are equal!
All parameters are equal!
All parameters are equal!
All parameters are equal!
All parameters are equal!

1.0.2 2. Cài đặt chiến lược thay đổi tốc độ học: Warm-up + Cosine Annealing LR

```
[ ]: def load_data(data_dir="./data"):  
    transform = transforms.Compose([  
        transforms.ToTensor(),  
        transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),  
    ])  
  
    trainset = torchvision.datasets.CIFAR10(  
        root=data_dir, train=True, download=True, transform=transform)  
  
    testset = torchvision.datasets.CIFAR10(  
        root=data_dir, train=False, download=True, transform=transform)  
  
    return trainset, testset  
  
class Net(nn.Module):  
    def __init__(self, l1=120, l2=84):  
        super(Net, self).__init__()  
        self.conv1 = nn.Conv2d(3, 6, 5)  
        self.pool = nn.MaxPool2d(2, 2)  
        self.conv2 = nn.Conv2d(6, 16, 5)  
        self.fc1 = nn.Linear(16 * 5 * 5, l1)  
        self.fc2 = nn.Linear(l1, l2)  
        self.fc3 = nn.Linear(l2, 10)  
        self.softmax = nn.Softmax()  
  
    def forward(self, x):  
        x = self.pool(F.relu(self.conv1(x)))  
        x = self.pool(F.relu(self.conv2(x)))  
        x = x.view(-1, 16 * 5 * 5)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        x = self.softmax(x)  
  
    return x
```

```
[ ]: trainset, testset = load_data('./data')
trainloader = torch.utils.data.DataLoader(
    trainset,
    batch_size=128,
    shuffle=True,
)

epochs = 50
warm_epoch = 5
init_lr = 1e-1
last_lr = 1e-5
T_max = epochs
T_cur = 0
lr_list = [0]

net = Net()
device = "cpu"
if torch.cuda.is_available():
    device = "cuda:0"
    if torch.cuda.device_count() > 1:
        net = nn.DataParallel(net)
net.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=init_lr, momentum=0.9)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to
./data/cifar-10-python.tar.gz

0%| | 0/170498071 [00:00<?, ?it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified

```
[ ]: for epoch in range(1, epochs+1): # loop over the dataset multiple times
    running_loss = 0.0
    epoch_steps = 0
    T_cur += 1

    # warm-up
    if epoch <= warm_epoch:
        optimizer.param_groups[0]['lr'] = (1.0 * epoch) / warm_epoch * init_lr
    else:
        # cosine annealing lr
        optimizer.param_groups[0]['lr'] = last_lr + (init_lr - last_lr) * (1 +
→np.cos(T_cur * np.pi / T_max)) / 2

    for i, data in enumerate(trainloader, 0):
```

```

    # get the inputs; data is a list of [inputs, labels]
    inputs, labels = data
    inputs, labels = inputs.to(device), labels.to(device)

    # zero the parameter gradients
    optimizer.zero_grad()

    # forward + backward + optimize
    outputs = net(inputs)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()

    # print statistics
    running_loss += loss.item()
    epoch_steps += 1
    if i + 1 == len(trainloader):
        print("[Epoch %d] loss: %.3f" % (epoch, running_loss / epoch_steps))
        running_loss = 0.0

    lr_list.append(optimizer.param_groups[0]['lr'])

print("Finished Training")

```

/usr/local/lib/python3.7/dist-packages/torch/nn/functional.py:718: UserWarning: Named tensors and all their associated APIs are an experimental feature and subject to change. Please do not use them for anything important until they are released as stable. (Triggered internally at /pytorch/c10/core/TensorImpl.h:1156.)

return torch.max_pool2d(input, kernel_size, stride, padding, dilation, ceil_mode)

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:33: UserWarning: Implicit dimension choice for softmax has been deprecated. Change the call to include dim=X as an argument.

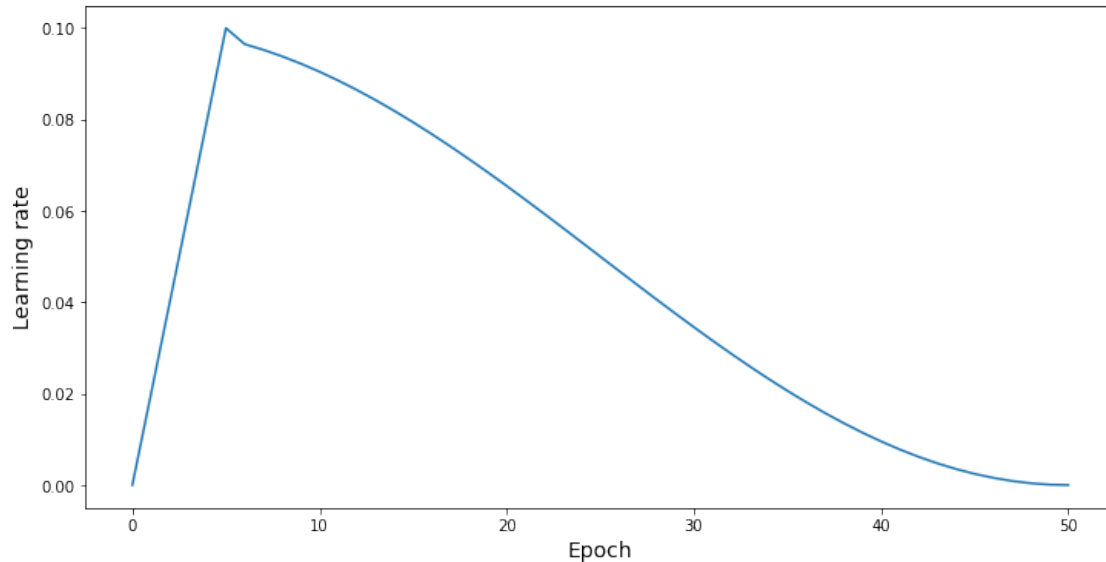
```

[Epoch 1] loss: 2.302
[Epoch 2] loss: 2.199
[Epoch 3] loss: 2.104
[Epoch 4] loss: 2.066
[Epoch 5] loss: 2.245
[Epoch 6] loss: 2.357
[Epoch 7] loss: 2.360
[Epoch 8] loss: 2.350
[Epoch 9] loss: 2.343
[Epoch 10] loss: 2.361
[Epoch 11] loss: 2.361
[Epoch 12] loss: 2.361
[Epoch 13] loss: 2.361

```

```
[Epoch 14] loss: 2.361
[Epoch 15] loss: 2.361
[Epoch 16] loss: 2.361
[Epoch 17] loss: 2.361
[Epoch 18] loss: 2.361
[Epoch 19] loss: 2.361
[Epoch 20] loss: 2.361
[Epoch 21] loss: 2.361
[Epoch 22] loss: 2.361
[Epoch 23] loss: 2.361
[Epoch 24] loss: 2.361
[Epoch 25] loss: 2.361
[Epoch 26] loss: 2.361
[Epoch 27] loss: 2.361
[Epoch 28] loss: 2.361
[Epoch 29] loss: 2.361
[Epoch 30] loss: 2.361
[Epoch 31] loss: 2.361
[Epoch 32] loss: 2.361
[Epoch 33] loss: 2.361
[Epoch 34] loss: 2.361
[Epoch 35] loss: 2.361
[Epoch 36] loss: 2.361
[Epoch 37] loss: 2.361
[Epoch 38] loss: 2.361
[Epoch 39] loss: 2.361
[Epoch 40] loss: 2.361
[Epoch 41] loss: 2.361
[Epoch 42] loss: 2.361
[Epoch 43] loss: 2.361
[Epoch 44] loss: 2.361
[Epoch 45] loss: 2.361
[Epoch 46] loss: 2.361
[Epoch 47] loss: 2.361
[Epoch 48] loss: 2.361
[Epoch 49] loss: 2.361
[Epoch 50] loss: 2.361
Finished Training
```

```
[ ]: plt.figure(figsize=(12, 6))
plt.plot(list(range(len(lr_list))), lr_list, label="lr")
plt.xlabel("Epoch", fontsize=14)
plt.ylabel("Learning rate", fontsize=14)
plt.show()
```

1.0.3 3. Tuning siêu tham số

Cài đặt thư viện ray

```
[ ]: !pip install ray
```

Collecting ray

Downloading ray-1.6.0-cp37-cp37m-manylinux2014_x86_64.whl (49.6 MB)

| 49.6 MB 6.2 kB/s

Requirement already satisfied: protobuf<=3.15.3 in /usr/local/lib/python3.7/dist-packages (from ray) (3.17.3)

Requirement already satisfied: numpy<=1.16 in /usr/local/lib/python3.7/dist-packages (from ray) (1.19.5)

Requirement already satisfied: filelock in /usr/local/lib/python3.7/dist-packages (from ray) (3.0.12)

Requirement already satisfied: pyyaml in /usr/local/lib/python3.7/dist-packages (from ray) (3.13)

Requirement already satisfied: msgpack<2.0.0,>=1.0.0 in /usr/local/lib/python3.7/dist-packages (from ray) (1.0.2)

Requirement already satisfied: grpcio<=1.28.1 in /usr/local/lib/python3.7/dist-packages (from ray) (1.39.0)

Requirement already satisfied: click<=7.0 in /usr/local/lib/python3.7/dist-packages (from ray) (7.1.2)

Collecting redis<=3.5.0

Downloading redis-3.5.3-py2.py3-none-any.whl (72 kB)

| 72 kB 565 kB/s

Requirement already satisfied: attrs in /usr/local/lib/python3.7/dist-packages (from ray) (21.2.0)

Requirement already satisfied: six<=1.5.2 in /usr/local/lib/python3.7/dist-

packages (from grpcio>=1.28.1->ray) (1.15.0)
Installing collected packages: redis, ray
Successfully installed ray-1.6.0 redis-3.5.3

```
[ ]: import os
      from ray.tune import CLIReporter
      from ray import tune
      from torch.utils.data import random_split
      from functools import partial
```

Viết hàm huấn luyện mô hình nơ-ron trên tập CIFAR10

```
[ ]: def train_cifar(config, checkpoint_dir=None, data_dir=None):
      net = Net(config["l1"], config["l2"])

      device = "cpu"
      if torch.cuda.is_available():
          device = "cuda:0"
          if torch.cuda.device_count() > 1:
              net = nn.DataParallel(net)
      net.to(device)

      criterion = nn.CrossEntropyLoss()
      optimizer = optim.SGD(net.parameters(), lr=config["lr"], momentum=0.9)

      if checkpoint_dir:
          model_state, optimizer_state = torch.load(
              os.path.join(checkpoint_dir, "checkpoint"))
          net.load_state_dict(model_state)
          optimizer.load_state_dict(optimizer_state)

      trainset, testset = load_data(data_dir)

      test_abs = int(len(trainset) * 0.8)
      train_subset, val_subset = random_split(
          trainset, [test_abs, len(trainset) - test_abs])

      trainloader = torch.utils.data.DataLoader(
          train_subset,
          batch_size=int(config["batch_size"]),
          shuffle=True,
          num_workers=8)
      valloader = torch.utils.data.DataLoader(
          val_subset,
          batch_size=int(config["batch_size"]),
          shuffle=True,
          num_workers=8)
```

```

for epoch in range(5): # loop over the dataset multiple times
    running_loss = 0.0
    epoch_steps = 0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        epoch_steps += 1
        if i % 2000 == 1999: # print every 2000 mini-batches
            print("[%d, %5d] loss: %.3f" % (epoch + 1, i + 1,
                                             running_loss / epoch_steps))

            running_loss = 0.0

    # Validation loss
    val_loss = 0.0
    val_steps = 0
    total = 0
    correct = 0
    for i, data in enumerate(valloader, 0):
        with torch.no_grad():
            inputs, labels = data
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = net(inputs)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

            loss = criterion(outputs, labels)
            val_loss += loss.cpu().numpy()
            val_steps += 1

    with tune.checkpoint_dir(epoch) as checkpoint_dir:
        path = os.path.join(checkpoint_dir, "checkpoint")
        torch.save((net.state_dict(), optimizer.state_dict()), path)

```

```

        print("*"*10, path, "*" * 10)

        tune.report(loss=(val_loss / val_steps), accuracy=correct / total)
    print("Finished Training")

```

Lựa chọn siêu tham số tốt (hyperparameter tuning)

```

[ ]: def main(num_samples=10, gpus_per_trial=2):
    data_dir = os.path.abspath("./data")
    config = {
        "l1": tune.grid_search([32, 64]),
        "l2": tune.grid_search([16, 32]),
        "lr": tune.grid_search([1e-4, 1e-2]),
        "batch_size": tune.grid_search([16])
    }
    reporter = CLIReporter(
        metric_columns=["loss", "accuracy", "training_iteration"])
    result = tune.run(
        partial(train_cifar, data_dir=data_dir),
        resources_per_trial={"cpu": 2, "gpu": gpus_per_trial},
        config=config,
        num_samples=num_samples,
        progress_reporter=reporter)

main(num_samples=10, gpus_per_trial=1)

```

```

/usr/local/lib/python3.7/dist-packages/ray/_private/services.py:238:
UserWarning: Not all Ray Dashboard dependencies were found. To use the dashboard
please install Ray using `pip install ray[default]`. To disable this message,
set RAY_DISABLE_IMPORT_WARNING env var to '1'.
    warnings.warn(warning_message)
2021-09-10 08:58:41,056 WARNING experiment.py:296 -- No name detected on
trainable. Using DEFAULT.
2021-09-10 08:58:41,058 INFO registry.py:67 -- Detected unknown callable for
trainable. Converting to class.
2021-09-10 08:58:41,115 INFO logger.py:597 -- pip install 'ray[tune]' to see
TensorBoard files.
2021-09-10 08:58:41,117 WARNING callback.py:117 -- The TensorboardX logger
cannot be instantiated because either TensorboardX or one of it's dependencies
is not installed. Please make sure you have the latest version of TensorboardX
installed: `pip install -U tensorboardx`

== Status ==
Memory usage on this node: 2.1/12.7 GiB
Using FIFO scheduling algorithm.
Resources requested: 2.0/2 CPUs, 1.0/1 GPUs, 0.0/7.32 GiB heap, 0.0/3.66 GiB
objects (0.0/1.0 accelerator_type:K80)
Result logdir: /root/ray_results/DEFAULT_2021-09-10_08-58-41
Number of trials: 16/80 (15 PENDING, 1 RUNNING)

```

Trial name	status	loc	batch_size	11	12	lr
DEFAULT_4b73f_00000	RUNNING		16	32	16	0.0001
DEFAULT_4b73f_00001	PENDING		16	64	16	0.0001
DEFAULT_4b73f_00002	PENDING		16	32	32	0.0001
DEFAULT_4b73f_00003	PENDING		16	64	32	0.0001
DEFAULT_4b73f_00004	PENDING		16	32	16	0.01
DEFAULT_4b73f_00005	PENDING		16	64	16	0.01
DEFAULT_4b73f_00006	PENDING		16	32	32	0.01
DEFAULT_4b73f_00007	PENDING		16	64	32	0.01
DEFAULT_4b73f_00008	PENDING		16	32	16	0.0001
DEFAULT_4b73f_00009	PENDING		16	64	16	0.0001
DEFAULT_4b73f_00010	PENDING		16	32	32	0.0001
DEFAULT_4b73f_00011	PENDING		16	64	32	0.0001
DEFAULT_4b73f_00012	PENDING		16	32	16	0.01
DEFAULT_4b73f_00013	PENDING		16	64	16	0.01
DEFAULT_4b73f_00014	PENDING		16	32	32	0.01
DEFAULT_4b73f_00015	PENDING		16	64	32	0.01

(pid=351) Files already downloaded and verified

(pid=351) Files already downloaded and verified

== Status ==

Memory usage on this node: 3.1/12.7 GiB

Using FIFO scheduling algorithm.

Resources requested: 2.0/2 CPUs, 1.0/1 GPUs, 0.0/7.32 GiB heap, 0.0/3.66 GiB objects (0.0/1.0 accelerator_type:K80)

Result logdir: /root/ray_results/DEFAULT_2021-09-10_08-58-41

Number of trials: 17/80 (16 PENDING, 1 RUNNING)

Trial name	status	loc	batch_size	11	12	lr
DEFAULT_4b73f_00000	RUNNING		16	32	16	0.0001
DEFAULT_4b73f_00001	PENDING		16	64	16	0.0001
DEFAULT_4b73f_00002	PENDING		16	32	32	0.0001
DEFAULT_4b73f_00003	PENDING		16	64	32	0.0001
DEFAULT_4b73f_00004	PENDING		16	32	16	0.01
DEFAULT_4b73f_00005	PENDING		16	64	16	0.01
DEFAULT_4b73f_00006	PENDING		16	32	32	0.01
DEFAULT_4b73f_00007	PENDING		16	64	32	0.01
DEFAULT_4b73f_00008	PENDING		16	32	16	0.0001
DEFAULT_4b73f_00009	PENDING		16	64	16	0.0001
DEFAULT_4b73f_00010	PENDING		16	32	32	0.0001
DEFAULT_4b73f_00011	PENDING		16	64	32	0.0001
DEFAULT_4b73f_00012	PENDING		16	32	16	0.01
DEFAULT_4b73f_00013	PENDING		16	64	16	0.01

	DEFAULT_4b73f_00014		PENDING				16		32		32		0.01	
	DEFAULT_4b73f_00015		PENDING				16		64		32		0.01	
	DEFAULT_4b73f_00016		PENDING				16		32		16		0.0001	

+-----+-----+-----+-----+-----+-----+-----+-----+

```
(pid=351) /usr/local/lib/python3.7/dist-
packages/torch/utils/data/dataloader.py:481: UserWarning: This DataLoader will
create 8 worker processes in total. Our suggested max number of worker in
current system is 2, which is smaller than what this DataLoader is going to
create. Please be aware that excessive worker creation might get DataLoader
running slow or even freeze, lower the worker number to avoid potential
slowness/freeze if necessary.
(pid=351)  cpuset_checked))
(pid=351) /usr/local/lib/python3.7/dist-
packages/torch/nn/functional.py:718: UserWarning: Named tensors and all their
associated APIs are an experimental feature and subject to change. Please do not
use them for anything important until they are released as stable. (Triggered
internally at /pytorch/c10/core/TensorImpl.h:1156.)
(pid=351)  return torch.max_pool2d(input, kernel_size, stride,
padding, dilation, ceil_mode)
(pid=351) /usr/local/lib/python3.7/dist-
packages/ray/workers/default_worker.py:33: UserWarning: Implicit dimension
choice for softmax has been deprecated. Change the call to include dim=X as an
argument.
(pid=351)  type=str,

(pid=351) [1, 2000] loss: 2.303
Result for DEFAULT_4b73f_00000:
  accuracy: 0.0999
  date: 2021-09-10_08-59-13
  done: false
  experiment_id: a590202bec2f4c1ea826538c56035a9f
  hostname: f7aececd2bb3
  iterations_since_restore: 1
  loss: 2.3026386421203613
  node_ip: 172.28.0.2
  pid: 351
  should_checkpoint: true
  time_since_restore: 30.793678045272827
  time_this_iter_s: 30.793678045272827
  time_total_s: 30.793678045272827
  timestamp: 1631264353
  timesteps_since_restore: 0
  training_iteration: 1
  trial_id: 4b73f_00000
```

== Status ==

Memory usage on this node: 3.1/12.7 GiB

Using FIFO scheduling algorithm.

Resources requested: 2.0/2 CPUs, 1.0/1 GPUs, 0.0/7.32 GiB heap, 0.0/3.66 GiB objects (0.0/1.0 accelerator_type:K80)

Result logdir: /root/ray_results/DEFAULT_2021-09-10_08-58-41

Number of trials: 17/80 (16 PENDING, 1 RUNNING)

+-----+-----+-----+-----+-----+-----+-----+						
-----+-----+-----+-----+-----+-----+-----+						
Trial name		status	loc	batch_size	l1	l2
lr	loss	accuracy	training_iteration			
-----+-----+-----+-----+-----+-----+-----+						
-----+-----+-----+-----+-----+-----+-----+						
DEFAULT_4b73f_00000		RUNNING	172.28.0.2:351		16	32 16
0.0001	2.30264	0.0999		1		
DEFAULT_4b73f_00001		PENDING			16	64 16
0.0001						
DEFAULT_4b73f_00002		PENDING			16	32 32
0.0001						
DEFAULT_4b73f_00003		PENDING			16	64 32
0.0001						
DEFAULT_4b73f_00004		PENDING			16	32 16
0.01						
DEFAULT_4b73f_00005		PENDING			16	64 16
0.01						
DEFAULT_4b73f_00006		PENDING			16	32 32
0.01						
DEFAULT_4b73f_00007		PENDING			16	64 32
0.01						
DEFAULT_4b73f_00008		PENDING			16	32 16
0.0001						
DEFAULT_4b73f_00009		PENDING			16	64 16
0.0001						
DEFAULT_4b73f_00010		PENDING			16	32 32
0.0001						
DEFAULT_4b73f_00011		PENDING			16	64 32
0.0001						
DEFAULT_4b73f_00012		PENDING			16	32 16
0.01						
DEFAULT_4b73f_00013		PENDING			16	64 16
0.01						
DEFAULT_4b73f_00014		PENDING			16	32 32
0.01						
DEFAULT_4b73f_00015		PENDING			16	64 32
0.01						
DEFAULT_4b73f_00016		PENDING			16	32 16
0.0001						
+-----+-----+-----+-----+-----+-----+-----+						
-----+-----+-----+-----+-----+-----+-----+						

```
(pid=351) ***** /root/ray_results/DEFAULT_2021-09-10_08-58-41/
DEFAULT_4b73f_00000_0_batch_size=16,l1=32,l2=16,lr=0.0001_2021-09-10_08-58-41/ch
eckpoint_000000/checkpoint *****
```

```
(pid=351) [2, 2000] loss: 2.303
```

```
Result for DEFAULT_4b73f_00000:
```

```
accuracy: 0.0999
date: 2021-09-10_08-59-39
done: false
experiment_id: a590202bec2f4c1ea826538c56035a9f
hostname: f7aececd2bb3
iterations_since_restore: 2
loss: 2.3026155532836916
node_ip: 172.28.0.2
pid: 351
should_checkpoint: true
time_since_restore: 57.259856939315796
time_this_iter_s: 26.46617889404297
time_total_s: 57.259856939315796
timestamp: 1631264379
timesteps_since_restore: 0
training_iteration: 2
trial_id: 4b73f_00000
```

```
== Status ==
```

```
Memory usage on this node: 3.1/12.7 GiB
```

```
Using FIFO scheduling algorithm.
```

```
Resources requested: 2.0/2 CPUs, 1.0/1 GPUs, 0.0/7.32 GiB heap, 0.0/3.66 GiB
```

```
objects (0.0/1.0 accelerator_type:K80)
```

```
Result logdir: /root/ray_results/DEFAULT_2021-09-10_08-58-41
```

```
Number of trials: 17/80 (16 PENDING, 1 RUNNING)
```

```
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| Trial name          | status | loc          | batch_size | l1 | l2 |
lr | loss | accuracy | training_iteration |
|-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
| DEFAULT_4b73f_00000 | RUNNING | 172.28.0.2:351 | 16 | 32 | 16 |
0.0001 | 2.30262 | 0.0999 | 2 |
| DEFAULT_4b73f_00001 | PENDING | | 16 | 64 | 16 |
0.0001 | | | |
| DEFAULT_4b73f_00002 | PENDING | | 16 | 32 | 32 |
0.0001 | | | |
| DEFAULT_4b73f_00003 | PENDING | | 16 | 64 | 32 |
0.0001 | | | |
| DEFAULT_4b73f_00004 | PENDING | | 16 | 32 | 16 |
0.01 | | | |
```


	DEFAULT_4b73f_00005		PENDING				16		64		16	
0.01												
	DEFAULT_4b73f_00006		PENDING				16		32		32	
0.01												
	DEFAULT_4b73f_00007		PENDING				16		64		32	
0.01												
	DEFAULT_4b73f_00008		PENDING				16		32		16	
0.0001												
	DEFAULT_4b73f_00009		PENDING				16		64		16	
0.0001												
	DEFAULT_4b73f_00010		PENDING				16		32		32	
0.0001												
	DEFAULT_4b73f_00011		PENDING				16		64		32	
0.0001												
	DEFAULT_4b73f_00012		PENDING				16		32		16	
0.01												
	DEFAULT_4b73f_00013		PENDING				16		64		16	
0.01												
	DEFAULT_4b73f_00014		PENDING				16		32		32	
0.01												
	DEFAULT_4b73f_00015		PENDING				16		64		32	
0.01												
	DEFAULT_4b73f_00016		PENDING				16		32		16	
0.0001												
+-----+-----+-----+-----+-----+-----+												
-----+-----+-----+-----+-----+												

```
(pid=351) ***** /root/ray_results/DEFAULT_2021-09-10_08-58-41/
DEFAULT_4b73f_00000_0_batch_size=16,l1=32,l2=16,lr=0.0001_2021-09-10_08-58-41/ch
eckpoint_000001/checkpoint *****
```

Cài đặt thử viện skorch

```
[ ]: !pip install skorch
import numpy as np
from sklearn.datasets import make_classification
from torch import nn

from skorch import NeuralNetClassifier
from sklearn.model_selection import GridSearchCV
```

Collecting skorch

Downloading skorch-0.10.0-py3-none-any.whl (128 kB)

| | 128 kB 5.3 MB/s

Requirement already satisfied: numpy>=1.13.3 in

/usr/local/lib/python3.7/dist-packages (from skorch) (1.19.5)

Requirement already satisfied: scipy>=1.1.0 in /usr/local/lib/python3.7/dist-packages (from skorch) (1.4.1)

Requirement already satisfied: tqdm>=4.14.0 in /usr/local/lib/python3.7/dist-packages (from skorch) (4.62.0)
Requirement already satisfied: tabulate>=0.7.7 in /usr/local/lib/python3.7/dist-packages (from skorch) (0.8.9)
Requirement already satisfied: scikit-learn>=0.19.1 in /usr/local/lib/python3.7/dist-packages (from skorch) (0.22.2.post1)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.19.1->skorch) (1.0.1)
Installing collected packages: skorch
Successfully installed skorch-0.10.0

Sử dụng thư viện skorch để huấn luyện

```
[ ]: trainset, testset = load_data('./data')
(X, y) = np.asarray(trainset.data[:]), np.asarray(trainset.targets[:])
X = X.reshape((-1, 3, 32, 32))
X = X.astype(np.float32)
y = y.astype(np.int64)

net = NeuralNetClassifier(
    Net,
    max_epochs=5,
    lr=0.01,
    # Shuffle training data on each epoch
    iterator_train__shuffle=True,
)

# training with default config
net.fit(X, y)
y_proba = net.predict_proba(X)
```

Files already downloaded and verified
Files already downloaded and verified

/usr/local/lib/python3.7/dist-packages/torch/nn/functional.py:718: UserWarning:
Named tensors and all their associated APIs are an experimental feature and
subject to change. Please do not use them for anything important until they are
released as stable. (Triggered internally at
/pytorch/c10/core/TensorImpl.h:1156.)

return torch.max_pool2d(input, kernel_size, stride, padding, dilation,
ceil_mode)

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:33: UserWarning:
Implicit dimension choice for softmax has been deprecated. Change the call to
include dim=X as an argument.

epoch	train_loss	valid_acc	valid_loss	dur
1	2.1907	0.2514	2.0782	
10.1372				
2	2.0612	0.2378	2.0725	9.8438

3	2.0083	0.2707	1.9760
9.8169			
4	1.9511	0.3011	1.9104
9.8356			
5	1.8987	0.3052	1.9103
9.8972			

Dùng skorch để lựa chọn siêu tham số

```
[ ]: # hyperparameter tuning
net.set_params(train_split=False, verbose=0)
params = {
    'lr': [1e-4, 1e-2],
    'module__l1': [32, 64],
    'module__l2': [16, 32],
}

gs = GridSearchCV(net, params, refit=False, cv=3, scoring='accuracy', verbose=2)

gs.fit(X, y)
print("best score: {:.3f}, best params: {}".format(gs.best_score_, gs.
↪best_params_))
```

Fitting 3 folds for each of 8 candidates, totalling 24 fits

[CV] lr=0.0001, module__l1=32, module__l2=16 ...

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
 /usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:33: UserWarning:
 Implicit dimension choice for softmax has been deprecated. Change the call to
 include dim=X as an argument.

[CV] ... lr=0.0001, module__l1=32, module__l2=16, total= 38.0s

[CV] lr=0.0001, module__l1=32, module__l2=16 ...

[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 38.0s remaining: 0.0s

[CV] ... lr=0.0001, module__l1=32, module__l2=16, total= 37.9s

[CV] lr=0.0001, module__l1=32, module__l2=16 ...

[CV] ... lr=0.0001, module__l1=32, module__l2=16, total= 37.8s

[CV] lr=0.0001, module__l1=32, module__l2=32 ...

[CV] ... lr=0.0001, module__l1=32, module__l2=32, total= 37.5s

[CV] lr=0.0001, module__l1=32, module__l2=32 ...

[CV] ... lr=0.0001, module__l1=32, module__l2=32, total= 37.6s

[CV] lr=0.0001, module__l1=32, module__l2=32 ...

[CV] ... lr=0.0001, module__l1=32, module__l2=32, total= 37.6s

[CV] lr=0.0001, module__l1=64, module__l2=16 ...

[CV] ... lr=0.0001, module__l1=64, module__l2=16, total= 37.4s

[CV] lr=0.0001, module__l1=64, module__l2=16 ...

[CV] ... lr=0.0001, module__l1=64, module__l2=16, total= 37.9s

[CV] lr=0.0001, module__l1=64, module__l2=16 ...

[CV] ... lr=0.0001, module__l1=64, module__l2=16, total= 38.1s

```

[CV] lr=0.0001, module__l1=64, module__l2=32 ...
[CV] ... lr=0.0001, module__l1=64, module__l2=32, total= 39.0s
[CV] lr=0.0001, module__l1=64, module__l2=32 ...
[CV] ... lr=0.0001, module__l1=64, module__l2=32, total= 38.5s
[CV] lr=0.0001, module__l1=64, module__l2=32 ...
[CV] ... lr=0.0001, module__l1=64, module__l2=32, total= 37.5s
[CV] lr=0.01, module__l1=32, module__l2=16 ...
[CV] ... lr=0.01, module__l1=32, module__l2=16, total= 36.1s
[CV] lr=0.01, module__l1=32, module__l2=16 ...
[CV] ... lr=0.01, module__l1=32, module__l2=16, total= 36.3s
[CV] lr=0.01, module__l1=32, module__l2=16 ...
[CV] ... lr=0.01, module__l1=32, module__l2=16, total= 36.3s
[CV] lr=0.01, module__l1=32, module__l2=32 ...
[CV] ... lr=0.01, module__l1=32, module__l2=32, total= 36.8s
[CV] lr=0.01, module__l1=32, module__l2=32 ...
[CV] ... lr=0.01, module__l1=32, module__l2=32, total= 37.0s
[CV] lr=0.01, module__l1=32, module__l2=32 ...
[CV] ... lr=0.01, module__l1=32, module__l2=32, total= 36.6s
[CV] lr=0.01, module__l1=64, module__l2=16 ...
[CV] ... lr=0.01, module__l1=64, module__l2=16, total= 36.9s
[CV] lr=0.01, module__l1=64, module__l2=16 ...
[CV] ... lr=0.01, module__l1=64, module__l2=16, total= 37.7s
[CV] lr=0.01, module__l1=64, module__l2=16 ...
[CV] ... lr=0.01, module__l1=64, module__l2=16, total= 38.1s
[CV] lr=0.01, module__l1=64, module__l2=32 ...
[CV] ... lr=0.01, module__l1=64, module__l2=32, total= 37.5s
[CV] lr=0.01, module__l1=64, module__l2=32 ...
[CV] ... lr=0.01, module__l1=64, module__l2=32, total= 36.9s
[CV] lr=0.01, module__l1=64, module__l2=32 ...
[CV] ... lr=0.01, module__l1=64, module__l2=32, total= 37.0s
best score: 0.287, best params: {'lr': 0.01, 'module__l1': 32, 'module__l2': 32}

[Parallel(n_jobs=1)]: Done 24 out of 24 | elapsed: 15.0min finished

```