

## ensemble\_pytorch\_answer

December 20, 2021

## 1 BTVN: Training Neural Networks (Tiếp)

Trong phần này các bạn sẽ làm quen với kỹ thuật model ensemble để tăng độ chính xác khi suy diễn

```
[ ]: !nvidia-smi
from google.colab import drive
drive.mount('/content/drive')

import torch
import torch.nn as nn
import torch.optim as optim
import numpy as np
import glob
import cv2
import torch.nn.functional as F
from torch.autograd import Variable
import os

import torchvision
import torchvision.transforms as transforms

from torch.nn import CrossEntropyLoss, Dropout, Softmax, Linear, Conv2d, LayerNorm
import matplotlib.pyplot as plt
from torchsummary import summary
```

Thu Sep 9 09:01:58 2021

+-----+-----+-----+-----+							
NVIDIA-SMI 470.63.01		Driver Version: 460.32.03		CUDA Version: 11.2			
+-----+-----+-----+-----+							
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile Uncorr. ECC		
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	
						MIG M.	
=====+=====+=====+=====							
0	Tesla K80	Off	00000000:00:04.0	Off		0	
N/A	33C	P8	29W / 149W	0MiB / 11441MiB	0%	Default	
						N/A	

+-----+-----+-----+							
+-----+-----+-----+							
Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory	
	ID	ID				Usage	
=====							
No running processes found							
+-----+-----+-----+							

Mounted at /content/drive

Tải dữ liệu và cài đặt một kiến trúc mạng nơ-ron đơn giản theo mô tả phía dưới

```
[ ]: def load_data(data_dir="./data"):
    transform = transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
    ])

    trainset = torchvision.datasets.CIFAR10(
        root=data_dir, train=True, download=True, transform=transform)

    testset = torchvision.datasets.CIFAR10(
        root=data_dir, train=False, download=True, transform=transform)

    return trainset, testset

class Net(nn.Module):
    def __init__(self, l1=120, l2=84):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, l1)
        self.fc2 = nn.Linear(l1, l2)
        self.fc3 = nn.Linear(l2, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)

        return x
```

```

model = Net()
if torch.cuda.is_available():
    model.cuda()
summary(model, (3, 32, 32))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 6, 28, 28]	456
MaxPool2d-2	[-1, 6, 14, 14]	0
Conv2d-3	[-1, 16, 10, 10]	2,416
MaxPool2d-4	[-1, 16, 5, 5]	0
Linear-5	[-1, 120]	48,120
Linear-6	[-1, 84]	10,164
Linear-7	[-1, 10]	850

Total params: 62,006

Trainable params: 62,006

Non-trainable params: 0

Input size (MB): 0.01

Forward/backward pass size (MB): 0.06

Params size (MB): 0.24

Estimated Total Size (MB): 0.31

/usr/local/lib/python3.7/dist-packages/torch/nn/functional.py:718: UserWarning:  
Named tensors and all their associated APIs are an experimental feature and  
subject to change. Please do not use them for anything important until they are  
released as stable. (Triggered internally at  
/pytorch/c10/core/TensorImpl.h:1156.)

```

    return torch.max_pool2d(input, kernel_size, stride, padding, dilation,
    ceil_mode)

```

Hàm đánh giá độ chính xác trên tập test

```

[ ]: def test_accuracy(net, device="cpu"):
    correct = 0
    total = 0
    with torch.no_grad():
        for data in testloader:
            images, labels = data
            images, labels = images.to(device), labels.to(device)
            outputs = net(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

```

```
return correct / total
```

Hàm huấn luyện mô hình

```
[ ]: def train(net, criterion, optimizer, save_path, device="cpu"):
    T_cur = 0
    for epoch in range(1, epochs+1): # loop over the dataset multiple times
        running_loss = 0.0
        epoch_steps = 0
        T_cur += 1

        # warm-up
        if epoch <= warm_epoch:
            optimizer.param_groups[0]['lr'] = (1.0 * epoch) / warm_epoch *
→init_lr
        else:
            # cosine annealing lr
            optimizer.param_groups[0]['lr'] = last_lr + (init_lr - last_lr) *
→(1 + np.cos(T_cur * np.pi / T_max)) / 2

        for i, data in enumerate(trainloader, 0):
            # get the inputs; data is a list of [inputs, labels]
            inputs, labels = data
            inputs, labels = inputs.to(device), labels.to(device)

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = net(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # print statistics
            running_loss += loss.item()
            epoch_steps += 1
            if i + 1 == len(trainloader):
                print("[Epoch %d] loss: %.3f" % (epoch, running_loss /
→epoch_steps))
                running_loss = 0.0

    print("Finished Training")
    print("Test accuracy:", test_accuracy(net, device))
    torch.save(net.state_dict(), save_path)
```

Thiết lập các tham số và hai kiến trúc mạng khác nhau

```
[ ]: epochs = 10
      warm_epoch = 5
      init_lr = 1e-2
      last_lr = 1e-4
      T_max = epochs

      configs = [{'l1': 64, 'l2': 32}, {'l1': 128, 'l2': 64}]

      trainset, testset = load_data('./data')
      trainloader = torch.utils.data.DataLoader(
          trainset,
          batch_size=128,
          shuffle=True,
      )
      testloader = torch.utils.data.DataLoader(
          testset, batch_size=4, shuffle=False, num_workers=2)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to  
./data/cifar-10-python.tar.gz

0%| | 0/170498071 [00:00<?, ?it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data  
Files already downloaded and verified

Huấn luyện hai mạng mô tả trong configs

```
[ ]: os.makedirs('./snapshot', exist_ok=True)

      for i, cfg in enumerate(configs):
          print(cfg)
          net = Net(cfg['l1'], cfg['l2'])
          device = "cpu"
          if torch.cuda.is_available():
              device = "cuda:0"
              if torch.cuda.device_count() > 1:
                  net = nn.DataParallel(net)
          net.to(device)

          criterion = nn.CrossEntropyLoss()
          optimizer = optim.SGD(net.parameters(), lr=init_lr, momentum=0.9)

          save_path = f'./snapshot/model{i}.pth'
          train(net, criterion, optimizer, save_path, device)
```

```
{'l1': 64, 'l2': 32}
[Epoch 1] loss: 2.293
[Epoch 2] loss: 1.969
[Epoch 3] loss: 1.645
[Epoch 4] loss: 1.464
```

```

[Epoch 5] loss: 1.346
[Epoch 6] loss: 1.205
[Epoch 7] loss: 1.150
[Epoch 8] loss: 1.112
[Epoch 9] loss: 1.093
[Epoch 10] loss: 1.085
Finished Training
Test accuracy: 0.5958
{'l1': 128, 'l2': 64}
[Epoch 1] loss: 2.303
[Epoch 2] loss: 2.249
[Epoch 3] loss: 1.819
[Epoch 4] loss: 1.558
[Epoch 5] loss: 1.394
[Epoch 6] loss: 1.224
[Epoch 7] loss: 1.164
[Epoch 8] loss: 1.123
[Epoch 9] loss: 1.100
[Epoch 10] loss: 1.091
Finished Training
Test accuracy: 0.5889

```

Kết hợp kết quả hai mạng (ensemble)

```

[ ]: from tqdm import tqdm

def test_ensemble(device="cuda:0"):
    correct = 0
    total = 0
    with torch.no_grad():
        for data in tqdm(testloader):
            images, labels = data
            images, labels = images.to(device), labels.to(device)
            final_outputs = torch.zeros((4, 10)).to(device)
            for i, cfg in enumerate(configs):
                net = Net(cfg['l1'], cfg['l2'])
                net.to(device)
                net.load_state_dict(torch.load(f'./snapshot/model{i}.pth'))
                outputs = net(images)
                final_outputs = final_outputs.add(outputs)

            final_outputs.div(len(configs))
            _, predicted = torch.max(final_outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    return correct / total

```

```
[ ]: from tqdm import tqdm

def test_ensemble(device="cuda:0"):
    correct = 0
    total = 0
    with torch.no_grad():
        for data in tqdm(testloader):
            images, labels = data
            images, labels = images.to(device), labels.to(device)
            final_outputs = torch.zeros((4, 10)).to(device)
            for i, cfg in enumerate(configs):
                net = Net(cfg['l1'], cfg['l2'])
                net.to(device)
                net.load_state_dict(torch.load(f'./snapshot/model{i}.pth'))
                outputs = net(images)
                final_outputs = final_outputs.add(outputs)

            final_outputs.div(len(configs))
            _, predicted = torch.max(final_outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    return correct / total
```

```
[ ]: test_ensemble()
```

```
100%|      | 2500/2500 [00:37<00:00, 67.16it/s]
```

```
[ ]: 0.6165
```

```
[ ]:
```